

# Red Wine Quality Prediction

IMPORT LIBRARIES AS WELL AS DATASET

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import os
%matplotlib inline
warnings.filterwarnings('ignore')
```


```
In [2]: df = pd.read_csv("Downloads/Python_Project_6_SVM.csv")
```

```
In [3]: df
```

Out[3]:

|      | Quality_Category | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | total: |
|------|------------------|-----------------|------------|---------------|-----------|-------------------|--------|
| 0    | 0                | 0.30            | 0.34       | 1.6           | 0.049     | 14                |        |
| 1    | 0                | 0.23            | 0.32       | 8.5           | 0.058     | 47                |        |
| 2    | 0                | 0.28            | 0.40       | 6.9           | 0.050     | 30                |        |
| 3    | 0                | 0.32            | 0.16       | 7.0           | 0.045     | 30                |        |
| 4    | 0                | 0.27            | 0.36       | 20.7          | 0.045     | 45                |        |
| ...  | ...              | ...             | ...        | ...           | ...       | ...               | ...    |
| 4889 | 0                | 0.21            | 0.29       | 1.6           | 0.039     | 24                |        |
| 4890 | 0                | 0.32            | 0.36       | 8.0           | 0.047     | 57                |        |
| 4891 | 0                | 0.24            | 0.19       | 1.2           | 0.041     | 30                |        |
| 4892 | 1                | 0.29            | 0.30       | 1.1           | 0.022     | 20                |        |
| 4893 | 0                | 0.21            | 0.38       | 0.8           | 0.020     | 22                |        |

4894 rows × 10 columns



```
In [4]: #Make a Copy of the Original dataset Which can help me in future
df1 = df.copy(deep=True)
df2 = df.copy(deep=True)
```

## DATA PREPROCESSING

```
In [5]: #checking for the missing value
df.isnull().sum()
```

```
Out[5]: Quality_Category      0
volatileacidity             0
citricacid                  0
residualsugar               0
chlorides                   0
freesulfurdioxide           0
totalsulfurdioxide           0
density                     0
sulphates                   0
alcohol                     0
dtype: int64
```

```
In [6]: df.info()
```


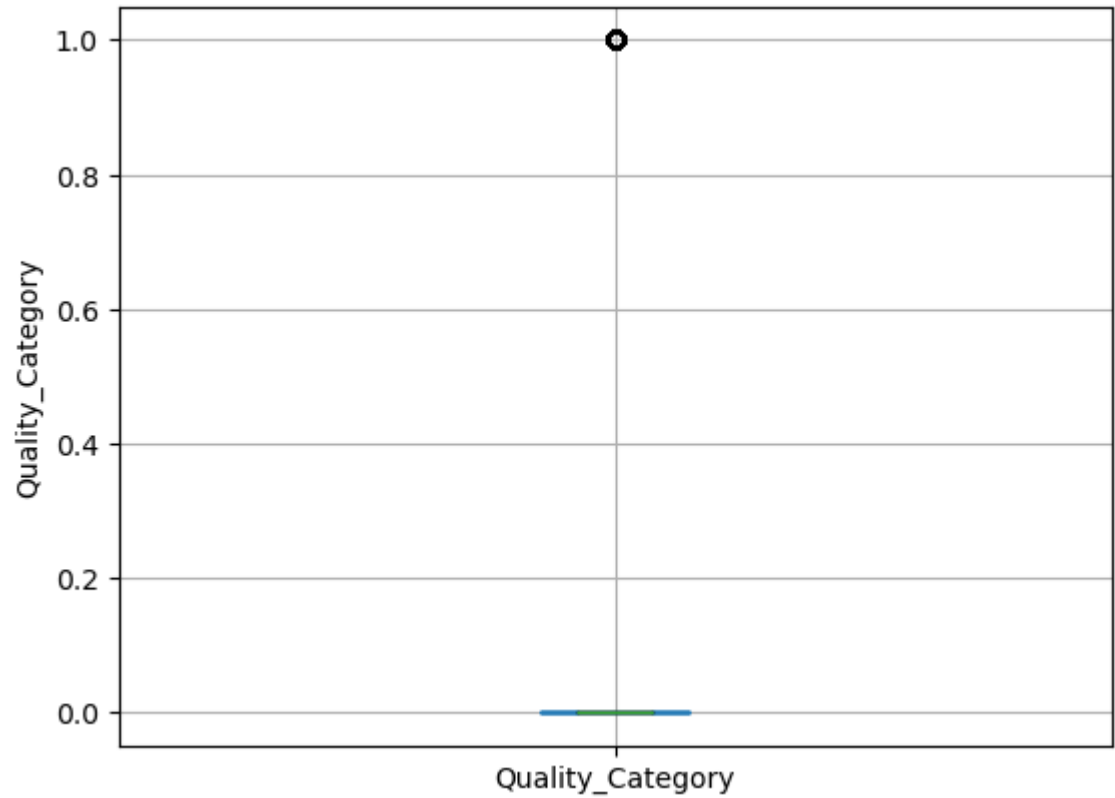
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4894 entries, 0 to 4893
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Quality_Category      4894 non-null   int64
1   volatileacidity        4894 non-null   float64
2   citricacid             4894 non-null   float64
3   residualsugar          4894 non-null   float64
4   chlorides              4894 non-null   float64
5   freesulfurdioxide      4894 non-null   int64
6   totalsulfurdioxide     4894 non-null   int64
7   density                4894 non-null   float64
8   sulphates              4894 non-null   float64
9   alcohol                4894 non-null   float64
dtypes: float64(7), int64(3)
memory usage: 382.5 KB
```


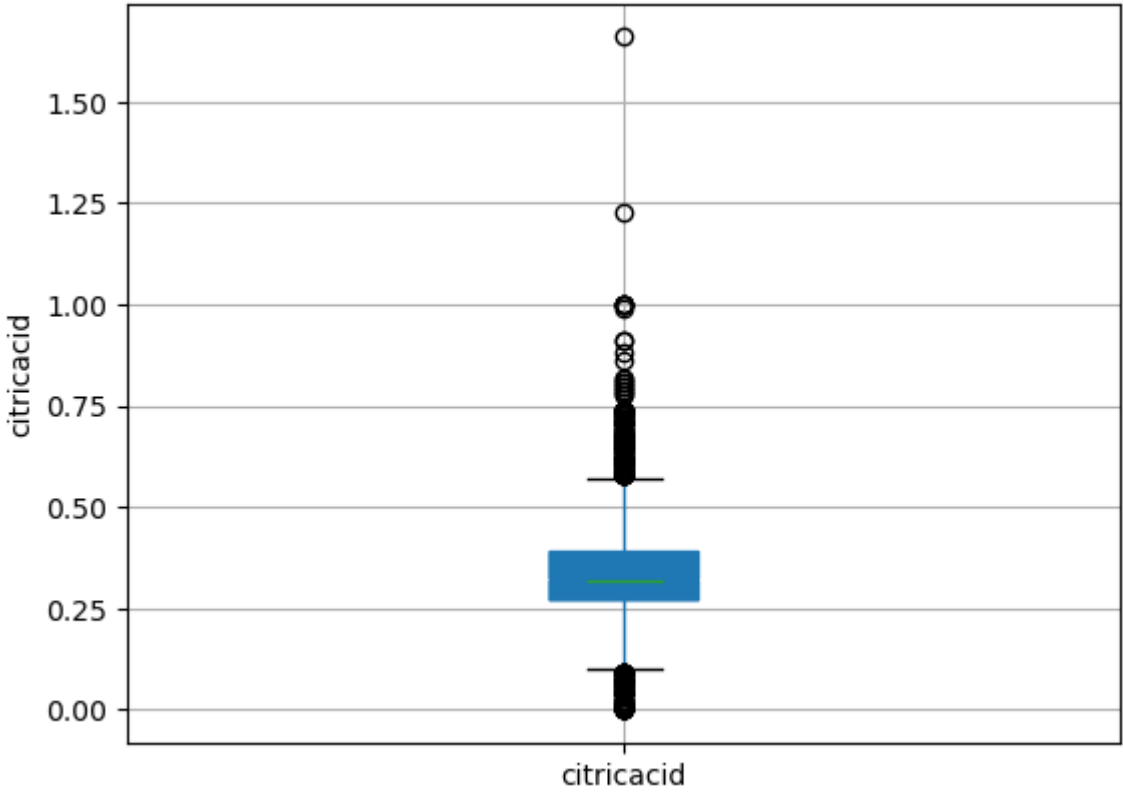
```
In [7]: df.columns
```


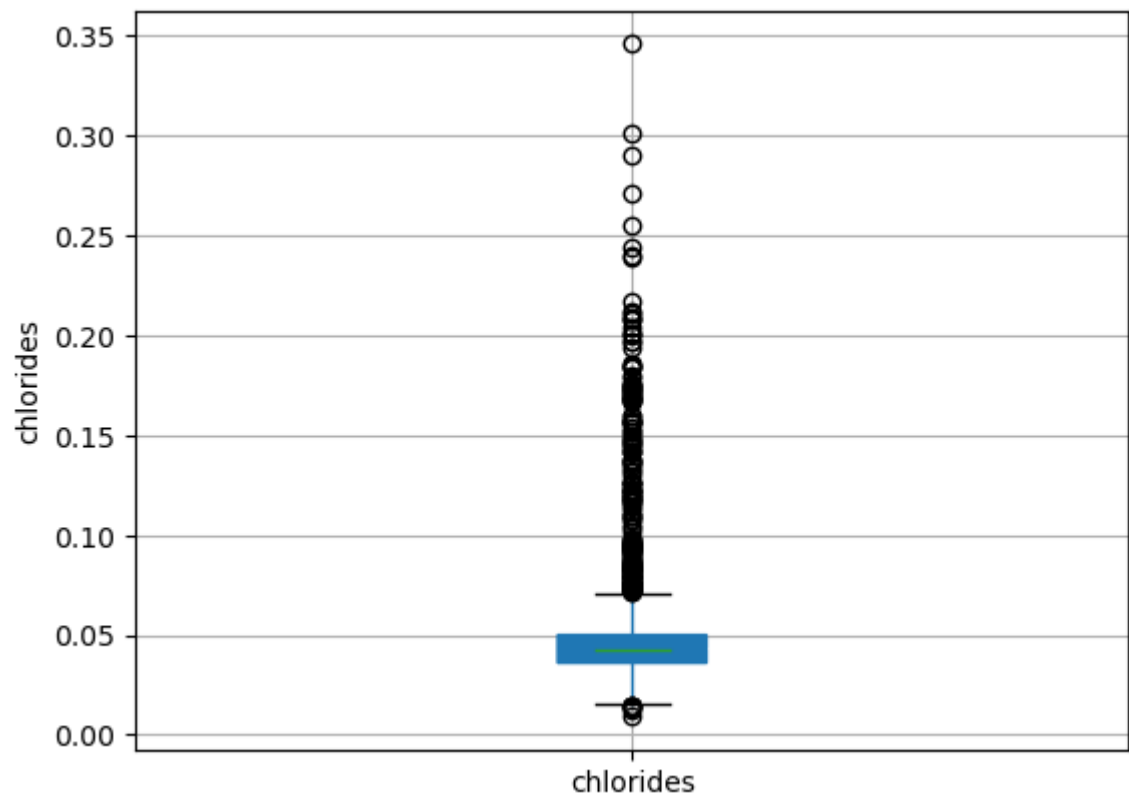
```
Out[7]: Index(['Quality_Category', 'volatileacidity', 'citricacid', 'residualsugar',
               'chlorides', 'freesulfurdioxide', 'totalsulfurdioxide', 'density',
               'sulphates', 'alcohol'],
              dtype='object')
```



## EXPLORATORY DATA ANALYSIS



```
In [8]: #Here, in this dataset have not Categorical features. So, now create EDA for numerical data.  
#First, check the Outliers  
for i in df:  
    df.boxplot(column=i, patch_artist = True, notch = 'True')  
    plt.ylabel(i)  
    plt.show()
```

















In [9]: *#we can see that in the numerical data has a outlier. So, we check distrubutio  
n of the numerical data*







```
In [10]: for i in df:
          sns.distplot(df[i], kde = True, color = 'green')
          plt.show()
```

















In [11]: *#we can see that in the dataset numerical distrubition in normal.*



```
In [12]: #Now, Lets check the correlation of the input and output Features.  
for i in df:  
    sns.scatterplot(df, y=df["Quality_Category"], x=df[i])  
    plt.show()
```











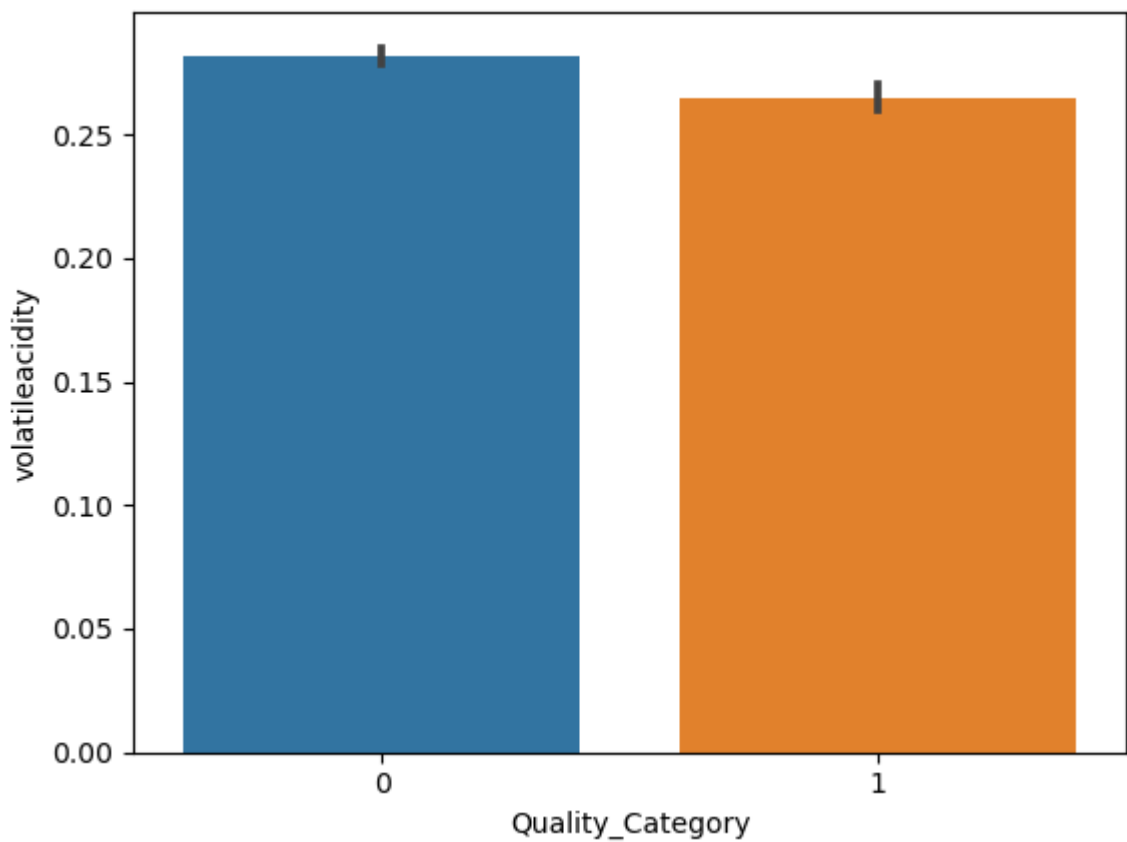







```
In [13]: df.columns
```



```
Out[13]: Index(['Quality_Category', 'volatileacidity', 'citricacid', 'residualsugar',  
               'chlorides', 'freesulfurdioxide', 'totalsulfurdioxide', 'density',  
               'sulphates', 'alcohol'],  
              dtype='object')
```



```
In [14]: for i in df:
          sns.barplot(df, x=df["Quality_Category"], y=df[i])
          plt.show()
```















In [15]: *#50, ALL the Features of data have important for Wine Quality.*

```
In [16]: #Check the duplicate
df[df.duplicated]
```

Out[16]:

|      | Quality_Category | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | total: |
|------|------------------|-----------------|------------|---------------|-----------|-------------------|--------|
| 5    | 0                | 0.30            | 0.34       | 1.6           | 0.049     | 14                |        |
| 35   | 0                | 0.24            | 0.39       | 18.0          | 0.057     | 45                |        |
| 44   | 0                | 0.31            | 0.26       | 7.4           | 0.069     | 28                |        |
| 57   | 0                | 0.19            | 0.26       | 12.4          | 0.048     | 50                |        |
| 59   | 0                | 0.38            | 0.15       | 4.6           | 0.044     | 25                |        |
| ...  | ...              | ...             | ...        | ...           | ...       | ...               | ...    |
| 4846 | 0                | 0.36            | 0.35       | 2.5           | 0.048     | 67                |        |
| 4847 | 0                | 0.33            | 0.44       | 8.9           | 0.055     | 52                |        |
| 4852 | 0                | 0.23            | 0.39       | 13.7          | 0.058     | 26                |        |
| 4876 | 0                | 0.34            | 0.40       | 8.1           | 0.046     | 68                |        |
| 4885 | 0                | 0.24            | 0.27       | 11.8          | 0.030     | 34                |        |

955 rows × 10 columns

Here, in this dataset have almost 955 row of duplicate but we can not delete because every red wine have their own different quality need that's why we don't drop duplicate.

## Seperate data in X and Y as well as Split data into train and Test

```
In [17]: # I am using a df1 data which was copy of the original data set.
x = df1.drop(["Quality_Category"], axis=1)
y = df1["Quality_Category"]
```


```
In [18]: from sklearn.model_selection import train_test_split
train_x, test_x, train_y, test_y = train_test_split(x,y, random_state=50, test_size=0.2, stratify=y)
```

In [19]: train\_x

Out[19]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | density |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|---------|
| 3073 | 0.25            | 0.38       | 7.9           | 0.045     | 54                | 208                | 0.994   |
| 2483 | 0.27            | 0.27       | 1.7           | 0.034     | 25                | 122                | 0.995   |
| 2155 | 0.24            | 0.37       | 1.8           | 0.031     | 6                 | 61                 | 0.994   |
| 4235 | 0.28            | 0.36       | 1.8           | 0.041     | 38                | 90                 | 0.994   |
| 635  | 0.34            | 0.14       | 5.8           | 0.046     | 49                | 197                | 0.994   |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...     |
| 3851 | 0.18            | 0.29       | 4.6           | 0.032     | 68                | 137                | 0.994   |
| 2015 | 0.32            | 0.22       | 16.7          | 0.046     | 38                | 133                | 0.994   |
| 3337 | 0.17            | 0.27       | 4.6           | 0.050     | 23                | 98                 | 0.994   |
| 1978 | 0.20            | 0.30       | 14.2          | 0.056     | 53                | 213                | 0.994   |
| 1475 | 0.16            | 0.49       | 2.0           | 0.056     | 20                | 124                | 0.994   |

3915 rows × 9 columns



```
In [20]: #create reset index
train_x.reset_index(inplace=True, drop=True)
test_x.reset_index(inplace=True, drop=True)

train_y.reset_index(inplace=True, drop=True)
test_y.reset_index(inplace=True, drop=True)
```

In [21]: train\_x

Out[21]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | dens |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 0    | 0.25            | 0.38       | 7.9           | 0.045     | 54                | 208                | 0.9  |
| 1    | 0.27            | 0.27       | 1.7           | 0.034     | 25                | 122                | 0.9  |
| 2    | 0.24            | 0.37       | 1.8           | 0.031     | 6                 | 61                 | 0.9  |
| 3    | 0.28            | 0.36       | 1.8           | 0.041     | 38                | 90                 | 0.9  |
| 4    | 0.34            | 0.14       | 5.8           | 0.046     | 49                | 197                | 0.9  |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...  |
| 3910 | 0.18            | 0.29       | 4.6           | 0.032     | 68                | 137                | 0.9  |
| 3911 | 0.32            | 0.22       | 16.7          | 0.046     | 38                | 133                | 0.9  |
| 3912 | 0.17            | 0.27       | 4.6           | 0.050     | 23                | 98                 | 0.9  |
| 3913 | 0.20            | 0.30       | 14.2          | 0.056     | 53                | 213                | 0.9  |
| 3914 | 0.16            | 0.49       | 2.0           | 0.056     | 20                | 124                | 0.9  |

3915 rows × 9 columns



In [22]: train\_y

Out[22]:

```
0      0
1      0
2      0
3      1
4      0
..
3910    0
3911    0
3912    0
3913    1
3914    0
```

Name: Quality\_Category, Length: 3915, dtype: int64

## Scaling Using Robustscaler

In [23]: `from sklearn.preprocessing import RobustScaler, MinMaxScaler, StandardScaler`  
`scaler = MinMaxScaler()`  
`scaler.fit(train_x)`

Out[23]:

```
▼ MinMaxScaler
MinMaxScaler()
```


In [24]: `train_x = pd.DataFrame(scaler.transform(train_x), columns=train_x.columns)`  
`test_x = pd.DataFrame(scaler.transform(test_x), columns=test_x.columns)`

In [25]: train\_x

Out[25]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | der  |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 0    | 0.166667        | 0.228916   | 0.111963      | 0.106825  | 0.181185          | 0.461717           | 0.17 |
| 1    | 0.186275        | 0.162651   | 0.016871      | 0.074184  | 0.080139          | 0.262181           | 0.07 |
| 2    | 0.156863        | 0.222892   | 0.018405      | 0.065282  | 0.013937          | 0.120650           | 0.05 |
| 3    | 0.196078        | 0.216867   | 0.018405      | 0.094955  | 0.125436          | 0.187935           | 0.05 |
| 4    | 0.254902        | 0.084337   | 0.079755      | 0.109792  | 0.163763          | 0.436195           | 0.13 |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...  |
| 3910 | 0.098039        | 0.174699   | 0.061350      | 0.068249  | 0.229965          | 0.296984           | 0.09 |
| 3911 | 0.235294        | 0.132530   | 0.246933      | 0.109792  | 0.125436          | 0.287703           | 0.21 |
| 3912 | 0.088235        | 0.162651   | 0.061350      | 0.121662  | 0.073171          | 0.206497           | 0.13 |
| 3913 | 0.117647        | 0.180723   | 0.208589      | 0.139466  | 0.177700          | 0.473318           | 0.23 |
| 3914 | 0.078431        | 0.295181   | 0.021472      | 0.139466  | 0.062718          | 0.266821           | 0.15 |

3915 rows × 9 columns



In [26]: train\_x.isnull().sum().sum()

Out[26]: 0

In [27]: test\_x.isnull().sum().sum()

Out[27]: 0

## Model Building And Evaluation

```
In [28]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
import xgboost as Xgb
```

```
In [29]: from sklearn.metrics import classification_report, accuracy_score, precision_s
core, recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay
```

```
In [30]: #LOGISTIC REGRESSION
log_model = LogisticRegression(random_state=50)
log_model.fit(train_x, train_y)
pred_log = log_model.predict(test_x)
print(classification_report(test_y, pred_log))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.95   | 0.88     | 767     |
| 1            | 0.55      | 0.22   | 0.32     | 212     |
| accuracy     |           |        | 0.79     | 979     |
| macro avg    | 0.68      | 0.59   | 0.60     | 979     |
| weighted avg | 0.76      | 0.79   | 0.76     | 979     |

```
In [31]: log_model.score(train_x, train_y)
```

```
Out[31]: 0.8068965517241379
```

```
In [32]: log_model.score(test_x, test_y)
```

```
Out[32]: 0.7916241062308478
```

```
In [33]: #KNEARASTNEIGHBORS CLASSIFIER
knn_model = KNeighborsClassifier(n_neighbors=10)
knn_model.fit(train_x, train_y)
pred_knn = knn_model.predict(test_x)
print(classification_report(test_y, pred_knn))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.85      | 0.92   | 0.89     | 767     |
| 1            | 0.61      | 0.42   | 0.50     | 212     |
| accuracy     |           |        | 0.82     | 979     |
| macro avg    | 0.73      | 0.67   | 0.69     | 979     |
| weighted avg | 0.80      | 0.82   | 0.80     | 979     |

```
In [34]: knn_model.score(train_x, train_y)
```

```
Out[34]: 0.8413793103448276
```

```
In [35]: knn_model.score(test_x, test_y)
```

```
Out[35]: 0.81511746680286
```



```
In [36]: # NAIVE BAYES CLASSIFICATION
nbc_model = GaussianNB()
nbc_model.fit(train_x, train_y)
pred_nbc = nbc_model.predict(test_x)
print(classification_report(test_y, pred_nbc))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.91      | 0.75   | 0.82     | 767     |
| 1            | 0.44      | 0.72   | 0.55     | 212     |
| accuracy     |           |        | 0.74     | 979     |
| macro avg    | 0.67      | 0.73   | 0.68     | 979     |
| weighted avg | 0.81      | 0.74   | 0.76     | 979     |

```
In [37]: nbc_model.score(train_x, train_y)
```

```
Out[37]: 0.7218390804597701
```

```
In [38]: nbc_model.score(test_x, test_y)
```

```
Out[38]: 0.7405515832482125
```

```
In [39]: # SUPPORT VECTOR CLASSIFICATION
svm_model = SVC(kernel="rbf")
svm_model.fit(train_x, train_y)
pred_svm = svm_model.predict(test_x)
print(classification_report(test_y, pred_svm))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.83      | 0.96   | 0.89     | 767     |
| 1            | 0.64      | 0.28   | 0.39     | 212     |
| accuracy     |           |        | 0.81     | 979     |
| macro avg    | 0.73      | 0.62   | 0.64     | 979     |
| weighted avg | 0.79      | 0.81   | 0.78     | 979     |

```
In [40]: svm_model.score(train_x, train_y)
```

```
Out[40]: 0.8209450830140486
```

```
In [41]: svm_model.score(test_x, test_y)
```

```
Out[41]: 0.8100102145045965
```

```
In [42]: #DECISION TREE CLASSIFICATION
dt_model = DecisionTreeClassifier(random_state=50, criterion="gini")
dt_model.fit(train_x, train_y)
pred_dt = dt_model.predict(test_x)
print(classification_report(test_y, pred_dt))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.88   | 0.89     | 767     |
| 1            | 0.61      | 0.66   | 0.63     | 212     |
| accuracy     |           |        | 0.83     | 979     |
| macro avg    | 0.76      | 0.77   | 0.76     | 979     |
| weighted avg | 0.84      | 0.83   | 0.84     | 979     |

```
In [43]: dt_model.score(train_x, train_y)
```

```
Out[43]: 1.0
```

```
In [44]: dt_model.score(test_x, test_y)
```

```
Out[44]: 0.8345250255362615
```

```
In [45]: #RANDOM FOREST CLASSIFICATION
rfc_model = RandomForestClassifier(random_state=50)
rfc_model.fit(train_x, train_y)
pred_rfc = rfc_model.predict(test_x)
print(classification_report(test_y, pred_rfc))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.96   | 0.93     | 767     |
| 1            | 0.80      | 0.61   | 0.70     | 212     |
| accuracy     |           |        | 0.88     | 979     |
| macro avg    | 0.85      | 0.79   | 0.81     | 979     |
| weighted avg | 0.88      | 0.88   | 0.88     | 979     |

```
In [46]: rfc_model.score(train_x, train_y)
```

```
Out[46]: 1.0
```

```
In [47]: rfc_model.score(test_x, test_y)
```

```
Out[47]: 0.8835546475995915
```

```
In [48]: #XGBOOST CLASSIFICATION
xgb_model = Xgb.XGBClassifier(n_estimators=100)
xgb_model.fit(train_x, train_y)
pred_xgb = xgb_model.predict(test_x)
print(classification_report(test_y, pred_xgb))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.95   | 0.93     | 767     |
| 1            | 0.77      | 0.63   | 0.69     | 212     |
| accuracy     |           |        | 0.88     | 979     |
| macro avg    | 0.84      | 0.79   | 0.81     | 979     |
| weighted avg | 0.87      | 0.88   | 0.87     | 979     |

```
In [49]: xgb_model.score(train_x, train_y)
```

```
Out[49]: 0.9961685823754789
```

```
In [50]: xgb_model.score(test_x, test_y)
```

```
Out[50]: 0.8794688457609806
```

```
In [51]: #ADABOOST CLASSIFICATION
from sklearn.ensemble import AdaBoostClassifier
adb_model = AdaBoostClassifier(random_state=50)
adb_model.fit(train_x, train_y)
pred_adb = adb_model.predict(test_x)
print(classification_report(test_y, pred_adb))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.84      | 0.93   | 0.89     | 767     |
| 1            | 0.60      | 0.36   | 0.45     | 212     |
| accuracy     |           |        | 0.81     | 979     |
| macro avg    | 0.72      | 0.65   | 0.67     | 979     |
| weighted avg | 0.79      | 0.81   | 0.79     | 979     |

```
In [52]: adb_model.score(train_x, train_y)
```

```
Out[52]: 0.8189016602809707
```

```
In [53]: adb_model.score(test_x, test_y)
```

```
Out[53]: 0.8100102145045965
```

Here, Decision Tree, Random forest and Xgboost model are overfitting of train and test dataset. So, we do a Hyper parameter tuning and Features selections.

# HYPERPARAMETER TUNING

```
In [54]: #HYPERPERAMETER TUNING OF LOGISTIC REGRESSOR
from sklearn.model_selection import GridSearchCV
log = LogisticRegression()
params = { "tol" : [0.1,0.5,0.8,0.9], "C" : [1,2,8,6,9],
           "solver": ['lbfgs', "liblinear", "newton-cg", "newton-cholesky", "saga", "saga"]}
clf1 = GridSearchCV(log, params, cv=5, scoring="accuracy")
clf1.fit(train_x, train_y)
print(clf1.best_params_)
print(clf1.best_score_)

{'C': 1, 'solver': 'newton-cholesky', 'tol': 0.1}
0.8066411238825031
```

```
In [55]: #HYPERPERAMETER TUNING OF KNN
knn = KNeighborsClassifier()
params_knn = {'algorithm' : ['auto', 'ball_tree', 'kd_tree', 'brute'], 'weights': ['uniform', 'distance'],
              "n_neighbors" : [1,25,14,13,26,85,45]}
clf2 = GridSearchCV(knn, params_knn, cv=5, scoring="accuracy")
clf2.fit(train_x, train_y)
print(clf2.best_params_)
print(clf2.best_score_)

{'algorithm': 'auto', 'n_neighbors': 45, 'weights': 'distance'}
0.8679438058748403
```

```
In [56]: #HYPERPERAMETER TUNING OF NB
nb = GaussianNB()
params_nb = {'var_smoothing' : [0.96,0.25,0.30,0.40, 0.50]}
clf3 = GridSearchCV(nb, params_nb, cv=5, scoring="accuracy")
clf3.fit(train_x, train_y)
print(clf3.best_params_)
print(clf3.best_score_)

{'var_smoothing': 0.5}
0.7984674329501915
```

```
In [57]: #HYPERPERAMETER TUNING OF SUPPORT VECTOR
svm = SVC()
params_svm = {"gamma" : ["scale", "auto"]}
clf4 = GridSearchCV(svm, params_svm, cv=5, scoring="accuracy")
clf4.fit(train_x, train_y)
print(clf4.best_params_)
print(clf4.best_score_)

{'gamma': 'scale'}
0.8153256704980842
```

```
In [58]: #HYPERPERAMETER TUNING OF DECISION TREE
dt = DecisionTreeClassifier()
params_dt = {'criterion':['gini', 'entropy', 'log_loss'], 'max_depth' : [1,25,14,13,45,75,26], 'splitter':['best', 'random']}
clf5 = GridSearchCV(dt, params_dt, cv=5, scoring="accuracy")
clf5.fit(train_x, train_y)
print(clf5.best_params_)
print(clf5.best_score_)
```

```
{'criterion': 'log_loss', 'max_depth': 75, 'splitter': 'best'}
0.823499361430396
```

```
In [59]: #HYPERPERAMETER TUNING OF RANDOMFOREST
rfc = RandomForestClassifier()
params_rfc = {"n_estimators" : [10,15,125,10,8,85], "max_depth" : [10,25,48,85,42,3]}
clf6 = GridSearchCV(rfc, params_rfc, cv=5, scoring="accuracy")
clf6.fit(train_x, train_y)
print(clf6.best_params_)
print(clf6.best_score_)
```

```
{'max_depth': 48, 'n_estimators': 125}
0.8720306513409962
```

```
In [60]: #HYPERPERAMETER TUNING OF XGBOOST
xgb = Xgb.XGBClassifier()
params_xgb = {'eta': [0.1, 0.2, 0.3,0.4,0.5], 'n_estimators' : [10, 50, 100,12,15], 'max_depth': [3, 6, 9,14]}
clf7 = GridSearchCV(xgb, params_xgb, cv=5, scoring="accuracy")
clf7.fit(train_x, train_y)
print(clf7.best_params_)
print(clf7.best_score_)
```

```
{'eta': 0.2, 'max_depth': 9, 'n_estimators': 100}
0.867432950191571
```

```
In [61]: #HYPERPERAMETER TUNING OF ADABOOST
adb = AdaBoostClassifier()
params_adb = {'n_estimators' : [10, 50, 100,12,15]}
clf8 = GridSearchCV(xgb, params_adb, cv=5, scoring="accuracy")
clf8.fit(train_x, train_y)
print(clf8.best_params_)
print(clf8.best_score_)
```

```
{'n_estimators': 100}
0.8592592592592592
```

```
In [62]: #best parameter for model
print("LogisticRegression score is :", clf1.best_params_)
print("KNeighborsClassifier score is :", clf2.best_params_)
print("GaussianNB score is :", clf3.best_params_)
print("Support vector machine score is :", clf4.best_params_)
print("DecisionTreeClassifier score is :", clf5.best_params_)
print("RandomForestClassifier score is :", clf6.best_params_)
print("XGB00ST score is :", clf7.best_params_)
print("AdaBoostClassifier score is :", clf8.best_params_)
```

```
LogisticRegression score is : {'C': 1, 'solver': 'newton-cholesky', 'tol': 0.1}
KNeighborsClassifier score is : {'algorithm': 'auto', 'n_neighbors': 45, 'weights': 'distance'}
GaussianNB score is : {'var_smoothing': 0.5}
Support vector machine score is : {'gamma': 'scale'}
DecisionTreeClassifier score is : {'criterion': 'log_loss', 'max_depth': 75, 'splitter': 'best'}
RandomForestClassifier score is : {'max_depth': 48, 'n_estimators': 125}
XGB00ST score is : {'eta': 0.2, 'max_depth': 9, 'n_estimators': 100}
AdaBoostClassifier score is : {'n_estimators': 100}
```

```
In [63]: #Score for all model
print("LogisticRegression score is :", clf1.best_score_)
print("KNeighborsClassifier score is :", clf2.best_score_)
print("GaussianNB score is :", clf3.best_score_)
print("Support vector machine score is :", clf4.best_score_)
print("DecisionTreeClassifier score is :", clf5.best_score_)
print("RandomForestClassifier score is :", clf6.best_score_)
print("XGB00ST score is :", clf7.best_score_)
print("AdaBoostClassifier score is :", clf8.best_score_)
```

```
LogisticRegression score is : 0.8066411238825031
KNeighborsClassifier score is : 0.8679438058748403
GaussianNB score is : 0.7984674329501915
Support vector machine score is : 0.8153256704980842
DecisionTreeClassifier score is : 0.823499361430396
RandomForestClassifier score is : 0.8720306513409962
XGB00ST score is : 0.867432950191571
AdaBoostClassifier score is : 0.8592592592592592
```

## Feature Selection

```
In [64]: #Correlation
corr = train_x.corr()
corr.style.background_gradient(cmap='coolwarm')
```

```
Out[64]:
```

|                    | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfur |
|--------------------|-----------------|------------|---------------|-----------|-------------------|-------------|
| volatileacidity    | 1.000000        | -0.155631  | 0.063650      | 0.086718  | -0.097849         | 0.0         |
| citricacid         | -0.155631       | 1.000000   | 0.096165      | 0.118984  | 0.095232          | 0.0         |
| residualsugar      | 0.063650        | 0.096165   | 1.000000      | 0.094485  | 0.311085          | 0.0         |
| chlorides          | 0.086718        | 0.118984   | 0.094485      | 1.000000  | 0.096071          | 0.0         |
| freesulfurdioxide  | -0.097849       | 0.095232   | 0.311085      | 0.096071  | 1.000000          | 0.0         |
| totalsulfurdioxide | 0.094374        | 0.122249   | 0.403077      | 0.195789  | 0.621998          | 1.0         |
| density            | 0.032416        | 0.160393   | 0.840780      | 0.255892  | 0.307270          | 0.0         |
| sulphates          | -0.026503       | 0.080483   | -0.024425     | 0.025187  | 0.070016          | 0.0         |
| alcohol            | 0.052898        | -0.083037  | -0.460439     | -0.355091 | -0.259787         | -0.0        |

```
In [65]: def correlation(dataset, threshold):
col_corr = set()
corr_matrix = dataset.corr()
for i in range(len(corr_matrix.columns)):
    for j in range(i):
        if abs(corr_matrix.iloc[i,j]) > threshold:
            colname = corr_matrix.columns[i]
            col_corr.add(colname)
return col_corr
```

```
In [66]: corr_features = correlation(train_x, 0.7)
len(set(corr_features))
```

```
Out[66]: 2
```

```
In [67]: corr_features
```

```
Out[67]: {'alcohol', 'density'}
```

```
In [68]: #Apply SelectKbest class to extract top Features
from sklearn.feature_selection import SelectKBest, chi2
bestfeatures = SelectKBest(score_func=chi2, k=7)
fit = bestfeatures.fit(x,y)
```

```
In [69]: fit
```

```
Out[69]:
```

▼
SelectKBest

SelectKBest(k=7, score\_func=<function chi2 at 0x0000026C3E79B7F0>)

```
In [70]: dfscores = pd.DataFrame(fit.scores_)
```

```
In [71]: dfcolumns = pd.DataFrame(x.columns)
```

```
In [72]: features = pd.concat([dfcolumns, dfscores], axis=1)
features.columns = ["specs", "score"]
```

```
In [73]: features
```

Out[73]:

|   | specs              | score       |
|---|--------------------|-------------|
| 0 | volatileacidity    | 0.849358    |
| 1 | citricacid         | 0.274252    |
| 2 | residualsugar      | 266.889301  |
| 3 | chlorides          | 1.706955    |
| 4 | freesulfurdioxide  | 21.459592   |
| 5 | totalsulfurdioxide | 1668.727815 |
| 6 | density            | 0.003509    |
| 7 | sulphates          | 0.295022    |
| 8 | alcohol            | 104.224001  |

## Feature Importance

```
In [74]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
model.fit(train_x, train_y)
```

Out[74]:


▾ ExtraTreesClassifier  
 ExtraTreesClassifier()

```
In [75]: print(model.feature_importances_)

[0.10391241 0.09376059 0.1042028  0.10170526 0.10340185 0.09875117
 0.09661143 0.09858443 0.19907007]
```



```
In [76]: feat_importance = pd.Series(model.feature_importances_, index=x.columns)
         feat_importance.nlargest(9).plot(kind="barh")
         plt.show()
```



```
In [77]: fe_model = RandomForestClassifier(random_state=50)
         fe_model.fit(train_x, train_y)
```

```
Out[77]: Random Forest Classifier
          RandomForestClassifier(random_state=50)
```

```
In [78]: feature_scores = pd.Series(fe_model.feature_importances_, index=train_x.columns)
         feature_scores.sort_values(ascending=False)
```

```
In [79]: feature_scores
```

```
Out[79]: alcohol      0.188746
         residualsugar 0.113954
         totalsulfurdioxide 0.107905
         chlorides    0.107803
         freesulfurdioxide 0.107724
         volatileacidity 0.107276
         sulphates    0.097564
         citricacid   0.093311
         density      0.075716
         dtype: float64
```

After the all Feature Selection method use then we decide to drop a density column for best accuracy. So, start with second time generate Model

## Start Whole Process of Model training in second time

In [80]: `df1.head()`

Out[80]:

|   | Quality_Category | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulf |
|---|------------------|-----------------|------------|---------------|-----------|-------------------|-----------|
| 0 | 0                | 0.30            | 0.34       | 1.6           | 0.049     | 14                |           |
| 1 | 0                | 0.23            | 0.32       | 8.5           | 0.058     | 47                |           |
| 2 | 0                | 0.28            | 0.40       | 6.9           | 0.050     | 30                |           |
| 3 | 0                | 0.32            | 0.16       | 7.0           | 0.045     | 30                |           |
| 4 | 0                | 0.27            | 0.36       | 20.7          | 0.045     | 45                |           |

In [81]: `X = df1.drop(["Quality_Category", "density"], axis=1)`  
`Y = df1["Quality_Category"]`

In [82]: `train_x1, test_x1, train_y1, test_y1 = train_test_split(X,Y, random_state=50, test_size=0.2, stratify=y)`

In [83]: `train_x1`

Out[83]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulp |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 3073 | 0.25            | 0.38       | 7.9           | 0.045     | 54                | 208                |      |
| 2483 | 0.27            | 0.27       | 1.7           | 0.034     | 25                | 122                |      |
| 2155 | 0.24            | 0.37       | 1.8           | 0.031     | 6                 | 61                 |      |
| 4235 | 0.28            | 0.36       | 1.8           | 0.041     | 38                | 90                 |      |
| 635  | 0.34            | 0.14       | 5.8           | 0.046     | 49                | 197                |      |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...  |
| 3851 | 0.18            | 0.29       | 4.6           | 0.032     | 68                | 137                |      |
| 2015 | 0.32            | 0.22       | 16.7          | 0.046     | 38                | 133                |      |
| 3337 | 0.17            | 0.27       | 4.6           | 0.050     | 23                | 98                 |      |
| 1978 | 0.20            | 0.30       | 14.2          | 0.056     | 53                | 213                |      |
| 1475 | 0.16            | 0.49       | 2.0           | 0.056     | 20                | 124                |      |

3915 rows × 8 columns

```
In [84]: #create reset index
train_x1.reset_index(inplace=True, drop=True)
test_x1.reset_index(inplace=True, drop=True)

train_y1.reset_index(inplace=True, drop=True)
test_y1.reset_index(inplace=True, drop=True)
```

```
In [85]: scaler.fit(train_x1)
train_x1 = pd.DataFrame(scaler.transform(train_x1), columns=train_x1.columns)
test_x1 = pd.DataFrame(scaler.transform(test_x1), columns=test_x1.columns)
```

```
In [86]: train_x1
```

Out[86]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulp |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 0    | 0.166667        | 0.228916   | 0.111963      | 0.106825  | 0.181185          | 0.461717           | 0.27 |
| 1    | 0.186275        | 0.162651   | 0.016871      | 0.074184  | 0.080139          | 0.262181           | 0.33 |
| 2    | 0.156863        | 0.222892   | 0.018405      | 0.065282  | 0.013937          | 0.120650           | 0.13 |
| 3    | 0.196078        | 0.216867   | 0.018405      | 0.094955  | 0.125436          | 0.187935           | 0.88 |
| 4    | 0.254902        | 0.084337   | 0.079755      | 0.109792  | 0.163763          | 0.436195           | 0.56 |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...  |
| 3910 | 0.098039        | 0.174699   | 0.061350      | 0.068249  | 0.229965          | 0.296984           | 0.18 |
| 3911 | 0.235294        | 0.132530   | 0.246933      | 0.109792  | 0.125436          | 0.287703           | 0.52 |
| 3912 | 0.088235        | 0.162651   | 0.061350      | 0.121662  | 0.073171          | 0.206497           | 0.29 |
| 3913 | 0.117647        | 0.180723   | 0.208589      | 0.139466  | 0.177700          | 0.473318           | 0.27 |
| 3914 | 0.078431        | 0.295181   | 0.021472      | 0.139466  | 0.062718          | 0.266821           | 0.37 |

3915 rows × 8 columns



```
In [87]: #KNeighborsClassifier
knn = KNeighborsClassifier(algorithm="auto", n_neighbors=25, weights="distance")
knn.fit(train_x1, train_y1)
pred2 = knn.predict(test_x1)
print(classification_report(test_y1, pred2))
print(accuracy_score(test_y1, pred2))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.94   | 0.92     | 767     |
| 1            | 0.74      | 0.64   | 0.69     | 212     |
| accuracy     |           |        | 0.87     | 979     |
| macro avg    | 0.82      | 0.79   | 0.80     | 979     |
| weighted avg | 0.87      | 0.87   | 0.87     | 979     |

0.874361593462717

```
In [88]: #LogisticRegression
log = LogisticRegression(C=1, solver="liblinear", tol=0.5)
log.fit(train_x1, train_y1)
pred1 = log.predict(test_x1)
print(classification_report(test_y1,pred1))
print(accuracy_score(test_y1, pred1))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 1.00   | 0.88     | 767     |
| 1            | 0.00      | 0.00   | 0.00     | 212     |
| accuracy     |           |        | 0.78     | 979     |
| macro avg    | 0.39      | 0.50   | 0.44     | 979     |
| weighted avg | 0.61      | 0.78   | 0.69     | 979     |

0.7834525025536262

```
In [89]: #GaussianNB
nb = GaussianNB(var_smoothing=0.5)
nb.fit(train_x1, train_y1)
pred3 = nb.predict(test_x1)
print(classification_report(test_y1,pred3))
print(accuracy_score(test_y1, pred3))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.95   | 0.88     | 767     |
| 1            | 0.56      | 0.25   | 0.34     | 212     |
| accuracy     |           |        | 0.79     | 979     |
| macro avg    | 0.69      | 0.60   | 0.61     | 979     |
| weighted avg | 0.76      | 0.79   | 0.76     | 979     |

0.7946884576098059

```
In [90]: #SVC
svm = SVC(gamma="auto")
svm.fit(train_x1, train_y1)
pred4 = svm.predict(test_x1)
print(classification_report(test_y1,pred4))
print(accuracy_score(test_y1, pred4))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 1.00   | 0.88     | 767     |
| 1            | 0.00      | 0.00   | 0.00     | 212     |
| accuracy     |           |        | 0.78     | 979     |
| macro avg    | 0.39      | 0.50   | 0.44     | 979     |
| weighted avg | 0.61      | 0.78   | 0.69     | 979     |

0.7834525025536262

```
In [91]: #DecisionTreeClassifier
dt = DecisionTreeClassifier(criterion="log_loss", max_depth=75, splitter="random")
dt.fit(train_x1, train_y1)
pred5 = dt.predict(test_x1)
print(classification_report(test_y1, pred5))
print(accuracy_score(test_y1, pred5))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.85   | 0.87     | 767     |
| 1            | 0.55      | 0.67   | 0.60     | 212     |
| accuracy     |           |        | 0.81     | 979     |
| macro avg    | 0.73      | 0.76   | 0.74     | 979     |
| weighted avg | 0.83      | 0.81   | 0.82     | 979     |

0.8089887640449438

```
In [92]: dt.score(train_x1, train_y1)
```

Out[92]: 1.0

```
In [93]: #RandomForestClassifier
rfc = RandomForestClassifier(max_depth=48, n_estimators=125)
rfc.fit(train_x1, train_y1)
pred6 = rfc.predict(test_x1)
print(classification_report(test_y1, pred6))
print(accuracy_score(test_y1, pred6))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.96   | 0.93     | 767     |
| 1            | 0.81      | 0.61   | 0.70     | 212     |
| accuracy     |           |        | 0.88     | 979     |
| macro avg    | 0.85      | 0.79   | 0.81     | 979     |
| weighted avg | 0.88      | 0.88   | 0.88     | 979     |

0.8845760980592441

```
In [94]: rfc.score(train_x1, train_y1)
```

Out[94]: 1.0

```
In [95]: #XGBClassifier
xgb = Xgb.XGBClassifier(eta=0.2 ,max_depth=9 ,n_estimators= 100)
xgb.fit(train_x1, train_y1)
pred7 = xgb.predict(test_x1)
print(classification_report(test_y1,pred7))
print(accuracy_score(test_y1, pred7))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.90      | 0.94   | 0.92     | 767     |
| 1            | 0.74      | 0.64   | 0.69     | 212     |
| accuracy     |           |        | 0.87     | 979     |
| macro avg    | 0.82      | 0.79   | 0.80     | 979     |
| weighted avg | 0.87      | 0.87   | 0.87     | 979     |

0.8733401430030644

```
In [96]: xgb.score(train_x1, train_y1)
```

Out[96]: 1.0

```
In [97]: #AdaBoostClassifier
adb = AdaBoostClassifier(n_estimators= 100)
adb.fit(train_x1, train_y1)
pred8 = adb.predict(test_x1)
print(classification_report(test_y1,pred8))
print(accuracy_score(test_y1, pred8))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.84      | 0.93   | 0.88     | 767     |
| 1            | 0.57      | 0.33   | 0.42     | 212     |
| accuracy     |           |        | 0.80     | 979     |
| macro avg    | 0.70      | 0.63   | 0.65     | 979     |
| weighted avg | 0.78      | 0.80   | 0.78     | 979     |

0.8018386108273748

Score After Hyper Parameter tuning and Feature Selection of all models

```
In [98]: print('LogisticRegression score is ', accuracy_score(test_y1, pred1))
print('KNeighborsClassifier score is', accuracy_score(test_y1, pred2))
print('GaussianNB score is', accuracy_score(test_y1, pred3))
print('Support Vector Machine score is', accuracy_score(test_y1, pred4))
print('DecisionTreeClassifier score is', accuracy_score(test_y1, pred5))
print('RandomForestClassifier score is', accuracy_score(test_y1, pred6))
print('XGBClassifier score is', accuracy_score(test_y1, pred7))
print('AdaBoostClassifier score is', accuracy_score(test_y1, pred8))
```

```
LogisticRegression score is 0.7834525025536262
KNeighborsClassifier score is 0.874361593462717
GaussianNB score is 0.7946884576098059
Support Vector Machine score is 0.7834525025536262
DecisionTreeClassifier score is 0.8089887640449438
RandomForestClassifier score is 0.8845760980592441
XGBClassifier score is 0.8733401430030644
AdaBoostClassifier score is 0.8018386108273748
```

CONCLUSION :- IN ABOVE GENERATED MODEL IN RANDOM FOREST AND XGBOOST CLASSIFIER GIVE ACCURACY IN TRAIN SET IS 100% BUT TESTING SET ACCURACY GIVES 88.25% AND 87.33%, RESPECTIVELY. SO, IN THIS CASE MODEL PERFORMING OVERFITTING. SO THAT WE DO A OVER SAMPLING BECAUSE THE DATASET HAVE INBALANCED SO WE DO IT AND CHECK THE ACCURACY OF MODEL.

## OVER SAMPLING

```
In [99]: #for over sampling import libraries.
```

```
In [100]: from imblearn.over_sampling import SMOTE
```

```
In [101]: os = SMOTE(random_state=50)
xos, yos = os.fit_resample(X, Y)
```

```
In [102]: train_x11, test_x11, train_y11, test_y11 = train_test_split(xos,yos, random_state=50, test_size=0.2)
```

```
In [103]: scaler.fit(train_x11)
train_x11 = pd.DataFrame(scaler.transform(train_x11), columns=train_x11.columns)
test_x11 = pd.DataFrame(scaler.transform(test_x11), columns=test_x11.columns)
```

In [104]: train\_x11

Out[104]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulp |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 0    | 0.441176        | 0.096386   | 0.032209      | 0.080119  | 0.111498          | 0.276102           | 0.39 |
| 1    | 0.135439        | 0.218328   | 0.039572      | 0.090630  | 0.195122          | 0.385151           | 0.32 |
| 2    | 0.137255        | 0.301205   | 0.200920      | 0.118694  | 0.188153          | 0.417633           | 0.55 |
| 3    | 0.362745        | 0.162651   | 0.062883      | 0.077151  | 0.052265          | 0.164733           | 0.17 |
| 4    | 0.088235        | 0.192771   | 0.015337      | 0.130564  | 0.156794          | 0.327146           | 0.68 |
| ...  | ...             | ...        | ...           | ...       | ...               | ...                | ...  |
| 6131 | 0.106997        | 0.168675   | 0.011740      | 0.107487  | 0.069686          | 0.185615           | 0.37 |
| 6132 | 0.098039        | 0.168675   | 0.147239      | 0.089021  | 0.094077          | 0.245940           | 0.29 |
| 6133 | 0.254902        | 0.198795   | 0.139571      | 0.080119  | 0.153310          | 0.378190           | 0.22 |
| 6134 | 0.137951        | 0.173579   | 0.010166      | 0.052204  | 0.101045          | 0.238979           | 0.36 |
| 6135 | 0.323529        | 0.132530   | 0.102761      | 0.077151  | 0.073171          | 0.250580           | 0.27 |

6136 rows × 8 columns



```
In [106]: #XGBClassifier
xgb.fit(train_x11, train_y11)
pred_1x = xgb.predict(test_x11)
print(classification_report(test_y11,pred_1x))
print(accuracy_score(test_y11, pred_1x))
print(xgb.score(train_x11, train_y11))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.92      | 0.90   | 0.91     | 735     |
| 1            | 0.91      | 0.92   | 0.92     | 799     |
| accuracy     |           |        | 0.91     | 1534    |
| macro avg    | 0.91      | 0.91   | 0.91     | 1534    |
| weighted avg | 0.91      | 0.91   | 0.91     | 1534    |


0.9146023468057366

0.9998370273794003



```
In [107]: print('XGBClassifier of confusion_matrix is:')
print(ConfusionMatrixDisplay.from_predictions(test_y11, pred_1x))
```

XGBClassifier of confusion\_matrix is:  
 <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x0000026C36AEA500>




```
In [108]: #RandomForestClassifier
rfc.fit(train_x11, train_y11)
pred_2r = rfc.predict(test_x11)
print(classification_report(test_y11,pred_2r))
print(accuracy_score(test_y11, pred_2r))
print(rfc.score(train_x11, train_y11))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.92      | 0.89   | 0.90     | 735     |
| 1            | 0.90      | 0.93   | 0.91     | 799     |
| accuracy     |           |        | 0.91     | 1534    |
| macro avg    | 0.91      | 0.91   | 0.91     | 1534    |
| weighted avg | 0.91      | 0.91   | 0.91     | 1534    |

0.9067796610169492  
 1.0

```
In [109]: print('RandomForestClassifier of confusion_matrix is:')
print(ConfusionMatrixDisplay.from_predictions(test_y11, pred_2r))
```

RandomForestClassifier of confusion\_matrix is:  
 <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x0000026C3EB97BE0>



## Testing the new data for checking

```
In [110]: x.tail()
```

Out[110]:

|      | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | dens |
|------|-----------------|------------|---------------|-----------|-------------------|--------------------|------|
| 4889 | 0.21            | 0.29       | 1.6           | 0.039     | 24                | 92                 | 0.9  |
| 4890 | 0.32            | 0.36       | 8.0           | 0.047     | 57                | 168                | 0.9  |
| 4891 | 0.24            | 0.19       | 1.2           | 0.041     | 30                | 111                | 0.9  |
| 4892 | 0.29            | 0.30       | 1.1           | 0.022     | 20                | 110                | 0.9  |
| 4893 | 0.21            | 0.38       | 0.8           | 0.020     | 22                | 98                 | 0.9  |

```
In [111]: new_df = {"volatileacidity": 0.30, 'citricacid': 0.31, "residualsugar": 1.0, 'ch
lorides': 0.20,
                  'freesulfurdioxide': 19, 'totalsulfurdioxide': 105, 'sulphates' : 0.
37, "alcohol": 12.0}
```

In [112]: index =[0]

In [113]: new\_df = pd.DataFrame(new\_df, index=index)  
new\_df

Out[113]:

|   | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulphate |
|---|-----------------|------------|---------------|-----------|-------------------|--------------------|----------|
| 0 | 0.3             | 0.31       | 1.0           | 0.2       | 19                | 105                | 0.3      |

In [114]: new\_df = pd.DataFrame(scaler.transform(new\_df), columns=new\_df.columns)

In [115]: new\_df

Out[115]:

|   | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulphate |
|---|-----------------|------------|---------------|-----------|-------------------|--------------------|----------|
| 0 | 0.215686        | 0.186747   | 0.006135      | 0.566766  | 0.059233          | 0.222738           | 0.18987  |

In [116]: prediction = xgb.predict(new\_df)  
prediction

Out[116]: array([0])

In [117]: prediction2 = rfc.predict(new\_df)  
prediction2

Out[117]: array([0], dtype=int64)

In [118]: new\_df1 = {"volatileacidity": 0.29, 'citricacid': 0.30, "residualsugar": 1.1, 'c  
hlorides': 0.022,  
                  'freesulfurdioxide': 20, 'totalsulfurdioxide': 110, 'sulphates' : 0.  
38, "alcohol": 12.8}

In [119]: new\_df1 = pd.DataFrame(new\_df1, index=index)  
new\_df1

Out[119]:

|   | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulphate |
|---|-----------------|------------|---------------|-----------|-------------------|--------------------|----------|
| 0 | 0.29            | 0.3        | 1.1           | 0.022     | 20                | 110                | 0.3      |

In [120]: new\_df1 = pd.DataFrame(scaler.transform(new\_df1), columns=new\_df1.columns)  
new\_df1

Out[120]:

|   | volatileacidity | citricacid | residualsugar | chlorides | freesulfurdioxide | totalsulfurdioxide | sulphate |
|---|-----------------|------------|---------------|-----------|-------------------|--------------------|----------|
| 0 | 0.205882        | 0.180723   | 0.007669      | 0.038576  | 0.062718          | 0.234339           | 0.20255  |

```
In [121]: prediction3 = xgb.predict(new_df1)
prediction3
```

```
Out[121]: array([1])
```

```
In [122]: prediction4 = rfc.predict(new_df1)
prediction4
```

```
Out[122]: array([1], dtype=int64)
```

```
In [123]: print('RandomForestClassifier score is ', accuracy_score(test_y11, pred_2r))
print('Xgboost Classifier score is', accuracy_score(test_y11, pred_1x))
```

```
RandomForestClassifier score is  0.9067796610169492
Xgboost Classifier score is 0.9146023468057366
```

CONCLUSION : From the above all Different Model's Random Forest Classification and Xgboost classifier have generated the model with higher accuracy in both default model and Hyperparameter tuning. But, in both model train score is higher than testing score. So, After Over sampling the dataset Xgboost give higher score and score is : 0.9146023468057366